

AI CONSULTING CAPSTONE • ROUND 2

Assist

Three capabilities. One controlled decision process.

What works, what does not yet pay, and how to decide with evidence.

Round 1 set the direction. Round 2 built the evidence

KEPT

Mid-size EU financial services firm

Complaint triage remained the lead use case

The system proposes; a person decides

CHANGED


One implemented use case became three

A POC became a working MVP

Cost estimates became ROI, risk and compliance assessments

Decision now: a staged evidence programme — not deployment.

One decision spine controls all three



In scope? Confident enough to propose? Every quote, figure or value present in the evidence?

The system never routes, files or blocks anything by itself.

The system proposes. A person decides.

1 · THE WORK

Complaint triage

WHAT Complaints, as they arrive.

WHY 64 overlapping categories; two handlers route the same one differently — and the clock is 15 days.

Anomaly review

WHAT Account transactions, in a batch.

WHY Unusual patterns surface as complaints, not alerts. By then the money has moved.

Reporting assistance

WHAT The quarter's complaint figures.

WHY Built by hand from spreadsheets, five days a quarter, and every figure has to be right.

2 · WHAT MAKES THE PROPOSAL

AI

The AI reads the complaint and picks the team from a fixed list — and must quote the customer's own sentence as its reason.

RULES, THEN AI

Rules pick the case out of the ledger. The AI never decides what is suspicious — it only explains the case the rules raised.

COMPUTED, THEN AI

The figures are calculated first. The AI only writes the narrative around numbers it did not produce.

3 · WHAT IT MUST PROVE

The sentence it quoted must appear in the complaint, **word for word**.

Every number in the note must **match the transaction record**.

Every figure must already be **on the computed sheet**.

4 · WHO DECIDES

A person. Every time.

It arrives either as a proposal with its evidence, or as a hold naming what stopped it. Nothing arrives decided.

Accept
agrees with it

Send elsewhere
overrides it

Escalate
wants another pair of eyes

Whichever they choose is written down beside what the system proposed.

One system, not three

the same checks and the same record serve all three — a fourth capability reuses them rather than starting again

Every step is kept

proposal, checks and human decision stored together — which is what makes agreement and override rates measurable at all


Minutes, not days

500 complaints classified in about a minute, eight at a time


The POC proves the mechanism — not accuracy

Proves: the workflow can propose, cite evidence and stop safely.

Does not prove: client performance or value.



The system works; the bespoke economics do not



At 2,426 complaints a year, bespoke deployment breaks even in month 46. Approximately 3,100 are needed.

€9,918 of €17,326 is cash — ombudsman fees and fraud losses, not saved time.

Three risks decide whether a pilot is defensible

Value assumptions fail on real data

4×5

Measure before committing to deployment. Fixed phase exits preserve the right to stop.

Handlers rubber-stamp proposals

4×4

Acceptance above 97% is a failure signal, measured per handler.

Oversight costs more than it saves

3×4

Measure effort during the shadow pilot from day one.

Not high-risk. Real-data use still has gates

AI Act:

Triage and reporting are not high-risk.

Anomaly flagging requires counsel confirmation. AI literacy applies now.

GDPR:

The full complaint narrative reaches the provider.

Close the DPA, transfer assessment and zero-retention controls before using real data.

The programme earns each next step

PHASE • WHAT IT BUYS

FEE • EXIT

0 • Ground truth, baselines, legal closure

€5,600 • go/no-go on pilot

1 • Configuration + 60-day shadow run

€12,600 • go/no-go on deployment

2 • Integration, training and handover

€14,000 • only if gates pass

€18,200 reaches the deployment decision. Nobody acts on a suggestion in shadow.

Hard gates: $\geq 80\%$ accuracy • 75–97% acceptance • zero ungrounded outputs.

The MVP makes every proposal accountable

One record connects the proposal, its evidence, every check and the human decision.

☐	2026-06-24	2 d	On track	Credit card	Disputes and fraud	Proposed	
--------------------------	------------	-----	----------	-------------	--------------------	----------	--

Proposed 15d left

Route to Payments and collections

Queue: Charged fees or interest you didn't expect

passed all checks

Because the customer wrote:

"I was charged a late fee while I was still trying to obtain the promised refund and dispute resolution."

> The full complaint

Nothing has happened yet. A person decides.

If rerouting, the right team is

Account servicing

Note (optional)

why the proposal was wrong

Accepted — routed as proposed

Rerouted — send to a different team

Sent to a person for full review

Checks

- ✓ input
- ✓ call
- ✓ parse
- ✓ scope
- ✓ confidence
- ✓ grounding

confidence 80% · 1759 ms · gpt-4o-mini · ref cx_896b75e4a1982a0c

Proceed to evidence — not deployment

Approve **Phases 0 and 1 up to €18,200.**

Close the legal and data-transfer gates before real-data processing.

Retain fixed phase exits and the stated pilot thresholds.

Do not fund the €14,000 deployment phase unless the evidence and commercial structure support it.

The POC has not yet proved client performance

MEASURE	RESULT
Complainant labels · gpt-4o-mini (deployed) / gpt-4o	56.8% / 60.5%
Self-agreement on the same input · gpt-4o	88.0%
Synthetic anomaly precision	97.9%
Accuracy / anomaly precision on client data	Not measured

Phase 0 creates expert labels; the shadow pilot tests real-world usefulness.

UC-1 and UC-3 are not high-risk; UC-2 needs counsel

STEP	FINDING
1 • AI system under Art. 3(1)	Yes
2 • Prohibited practice under Art. 5	None identified
4 • Annex III high-risk use	UC-1/3: no • UC-2: counsel
Articles 50 and 4	Marking + AI literacy
The intended use and human decision boundary drive the classification.	

The full complaint narrative still leaves the bank

WHAT

WHERE IT GOES

Complaint identifier

No • replaced before the model call

Complaint narrative

Yes • sent to the provider in full

Evidence + pseudonymous reference

Stored in monitoring

Pseudonymisation protects the reference, not the content. Transfer controls are the gate.

Every pilot threshold must pass before deployment

MEASURE	THRESHOLD
Accuracy against expert labels	$\geq 80\%$
Triage-time reduction	$\geq 40\%$
Handler acceptance	75–97%
Oversight effort	≤ 3 days / quarter
Ungrounded outputs reaching a handler	Zero
Correct abstention	$\geq 95\%$
Anomaly precision on disposition	$\geq 30\%$