

AI CONSULTING CAPSTONE • ROUND 2

Assist

Complaint operations for a mid-size EU bank.

Ugo Ahukannah

Financial services • mid-size EU retail bank



What changed after Round 1

TAKEAWAY Same client, same use case. The scope widened.

ROUND 1		ROUND 2
Three proposed, one built	→	Three built
A workflow you could watch	→	An application you can use
AI Act work deferred	→	Classified, with reasoning

Decision: KEEP. No customer-facing chatbot. No credit scoring.



Three jobs, one shape

TAKEAWAY High volume, low complexity, unstructured text.

Routing

2,426 a year, 64 categories. Two handlers disagree.

Reporting

The regulatory return, built by hand under deadline.

Fraud

7.4% unauthorised. The money has already gone.

A decision nobody can see is one nobody can sign off.



One decision spine

TAKEAWAY Every capability is the same four steps.




Only the evidence changes — a quote, a figure, a value.

Nothing routes, files or blocks by itself.



The proof of concept

TAKEAWAY Nine steps, left to right. Green means it ran.





What it returns

TAKEAWAY A bespoke build does not pay back at your size.

	ROI AT 36 MONTHS
Bespoke build, as proposed	-31.0%
Productised across five firms	+16.9%
...with oversight in-house	+96.6%

Break-even month 69. €18,200 buys a decision, not a deployment.



The three that could derail this

TAKEAWAY The two biggest are not technical.

RISK	MITIGATION
Value assumptions wrong 4×5	The pilot measures it. Stop at €18,200
Handlers rubber-stamp 4×4	Acceptance above 97% is a failure signal
Oversight costs more 3×4	One spine, one review. Instrumented

Twelve risks scored in full. These are the top three.



Where this sits in law

TAKEAWAY Not high-risk — and the reason matters more than the label.

EU AI ACT

- Not Annex III — never creditworthiness
- Art. 50 transparency applies
- Art. 4 AI literacy applies today

GDPR

- You are controller and provider
- The narrative goes to the model in full
- Transfer work gates any pilot

It proposes, a person decides. That is what holds the tier.



How this would go

TAKEAWAY Shadow mode first. Nobody acts on it.

0

Discovery

€5,600 • 4 weeks

1

Pilot

€12,600 • 12 weeks

2

Deploy

€14,000 • 10 weeks

Greenlight: accuracy $\geq 80\%$, acceptance 75–97%, zero ungrounded.



The beginning of the real product

☐	2026-06-22	3 d	12 d	on track	Credit card	Originations	0.80	proposed	OK_PROPOSED
☐	2026-06-17	8 d	7 d	on track	Checking or savings account	Account servicing	0.90	proposed	OK_PROPOSED
☐	2026-06-17	8 d	7 d	on track	Credit card	Credit reporting	0.80	proposed	OK_PROPOSED
☐	2026-06-16	9 d	6 d	on track	Checking or savings account	Disputes and fraud	0.80	proposed	OK_PROPOSED
☐	2026-06-14	11 d	4 d	due soon	Vehicle loan or lease	Asset recovery	0.90	proposed	OK_PROPOSED

Proposed 15d left OK_PROPOSED

Route to Disputes and fraud

Queue: Problem with a lender or other company charging your account

passed all checks

Because the customer wrote:

“Chase froze the funds, closed my account, and referred the matter to its fraud department.”

> The full complaint

Nothing has happened yet. A person decides.

Accepted — routed as proposed

Rerouted — handler chose a different team

Sent to a person for full review

Checks

- ✓ input
- ✓ call
- ✓ parse
- ✓ scope
- ✓ confidence
- ✓ grounding

confidence 0.80 · 2769 ms · gpt-4o-mini · ref cx_db78f958aa4d3606

- What the POC only implied
 - A queue, already read.
 - A deadline column.
 - Every check it ran, by name.
 - Nothing moves until someone presses a button.



What I'm asking for

€18,200 to Phase 1.

Nobody acts on a single suggestion.

- It works — three capabilities, every decision checked.
- It is lawful — not high-risk, reasoning written down.
- It does not yet pay for itself. The pilot settles that.



Backup



Backup • Why 60.5% is not accuracy

TAKEAWAY The model is consistent. It disagrees with an untrained label.

MEASURE	RESULT
Agreement with CFPB labels	60.5%
Self-agreement, same input twice	88.0%
Fabricated quotes caught	4 of 60

Phase 0 creates the expert labels an accuracy claim would need.



Backup • The anomaly detector's 100%

TAKEAWAY It finds every planted anomaly. That is not a detection rate.

WHAT IT SHOWS

- Deterministic — same batch, same result
- All four patterns, not just the easy one
- 97.1% precision on 2,069 transactions

WHAT IT DOES NOT

- Thresholds and generator, same author
- So it measures self-agreement
- Only a pilot on real traffic settles it

Same trap as the 60.5%.



Backup • AI Act classification

TAKEAWAY Seven steps. Step 4 is the one that matters.

STEP	FINDING
1 • An AI system?	Yes — Art. 3(1)
2 • Prohibited, Art. 5?	No — one close call, flagged
4 • An Annex III use case?	No — never creditworthiness
6 • Art. 50 transparency	Drafts marked anyway



Backup • What actually leaves the building

TAKEAWAY Pseudonymisation protects the reference, not the content.

WHAT	WHERE IT GOES
The complaint narrative	To the model provider, in full (US)
The quoted sentence	Stored in monitoring
The customer reference	Pseudonymised at the boundary

So the transfer work is a gate before any pilot on real data.



Backup • What the pilot has to prove

TAKEAWAY Miss any of the first three and we do not proceed.

MEASURE	THRESHOLD
Team accuracy vs expert labels	$\geq 80\%$
Acceptance rate, per handler	75 – 97%
Ungrounded output reaching a handler	Zero
Correct abstention	$\geq 95\%$
Anomaly precision on disposition	$\geq 30\%$



Backup • Anomaly flagging

☐	2026-06-20	2 d	threshold_structuring	150.7×	€4935.00	5	held	REJECT_VALUE_MISMATCH
☐	2026-06-20	2 d	amount_spike	54.8×	€995.00	1	proposed	OK_PROPOSED

Proposed OK_PROPOSED

Amount spike

a single amount far above this account's own normal

Transactions	Total	Account median	Times normal
1	€990	€18	54.5×

Case note:

This candidate was raised due to a significant spike in the transaction amount, with a single transfer of EUR 990.0, which is far above the account's median transaction of EUR 18.15. The amount is approximately 54.5 times higher than what is typical for this account, indicating an unusual activity pattern.

Fastest check: Review the account's transaction history for any similar high-value transfers.

Nothing has happened yet. A person decides.

Escalated to an analyst

Dismissed — ordinary for this account

Needs more information

Checks

- ✓ input
- ✓ call
- ✓ parse
- ✓ scope
- ✓ confidence
- ✓ grounding

confidence 0.80 · 1377 ms · gpt-4o-mini · ref cx_65b1bc11dceaf3b

transfer · online_banking · IE

■ The division of labour

- A detector decides **WHAT** is unusual.
- The model only explains **WHY**.
- Ranked by that account's own normal — never by amount.
- Six of 56 notes stopped for an uncited figure.