

AI CONSULTING CAPSTONE • ROUND 2

Assist

Complaint operations for a mid-size EU bank.

Ugo Ahukannah

Financial services • mid-size EU retail bank

Complaint triage Reporting assistance Anomaly flagging

What changed after Round 1

ROUND 1

Three use cases proposed, one built

A workflow you could watch

AI Act work deferred

ROUND 2

All three built

An application you can use

Classified, with the reasoning

Decision: KEEP. No customer-facing chatbot. No credit scoring.

Three jobs, one shape

Routing

2,426 complaints a year, 64 categories. Two handlers pick differently.

Reporting

The regulatory return, assembled by hand, under deadline.

Fraud

7.4% are unauthorised transactions. The money has already gone.

High volume, low complexity, unstructured text.

One decision spine


Validate input → Model proposes → Guards check → Person decides

Only the evidence changes — a verbatim quote, a figure we computed, or the transaction's own values.

Nothing routes, files or blocks by itself.

The proof of concept

Nine steps, left to right. Green means it ran.



What it returns

ROI AT 36 MONTHS

Bespoke build, as proposed	-31.0%
Productised across five firms	+16.9%
...with oversight in-house	+96.6%

Break-even month 69. €18,200 buys a decision, not a deployment.

The three that could derail this

The value assumptions are wrong 4×5

The pilot measures exactly that. Fixed fee — you can stop at €18,200.

Handlers rubber-stamp the proposal 4×4

Acceptance above 97% is a failure signal, measured per handler.

Oversight costs more than it saves 3×4

One spine, so one review covers all three. Instrumented from day one.

Twelve risks scored in full — regulatory, technical, ethical, operational.

Where this sits in law

Not high-risk. Not an Annex III use case — it never touches creditworthiness.

You are the controller, and also the provider.

The narrative reaches the model in full — pseudonyms protect the reference, not the content.

Transfer work is a gate before any pilot on real data.

It proposes, a person decides. That is what holds the tier.

How this would go

PHASE	COST • DURATION
Phase 0 • Discovery, expert labelling	€5,600 • 4 weeks
Phase 1 • Pilot, 60-day shadow run	€12,600 • 12 weeks
Phase 2 • Full deployment	€14,000 • 10 weeks

Phase 1 runs in shadow — nobody acts on it, nothing changes for a customer.

Phase 2 only if accuracy $\geq 80\%$, acceptance 75–97%, zero ungrounded outputs.

The beginning of the real product

A queue already read. Every check by name. Nothing moves until someone presses a button.

☐	2026-06-22	3 d	12 d	on track	Credit card	Originations	0.80	proposed	OK_PROPOSED
☐	2026-06-17	8 d	7 d	on track	Checking or savings account	Account servicing	0.90	proposed	OK_PROPOSED
☐	2026-06-17	8 d	7 d	on track	Credit card	Credit reporting	0.80	proposed	OK_PROPOSED
☐	2026-06-16	9 d	6 d	on track	Checking or savings account	Disputes and fraud	0.80	proposed	OK_PROPOSED
☐	2026-06-14	11 d	4 d	due soon	Vehicle loan or lease	Asset recovery	0.90	proposed	OK_PROPOSED

Proposed 15d left OK_PROPOSED

Route to Disputes and fraud

Queue: Problem with a lender or other company charging your account

passed all checks

Because the customer wrote:

"Chase froze the funds, closed my account, and referred the matter to its fraud department."

> The full complaint

Nothing has happened yet. A person decides.

Accepted — routed as proposed

Rerouted — handler chose a different team

Sent to a person for full review

Checks

- ✓ input
- ✓ call
- ✓ parse
- ✓ scope
- ✓ confidence
- ✓ grounding

confidence 0.80 · 2769 ms · gpt-4o-mini · ref cx_db78f958aa4d3606

What I'm asking for

€18,200 to Phase 1. Nobody acts on a single suggestion.

It works — three capabilities running, every decision checked.

It is lawful — not high-risk, with the reasoning written down.

It does not yet pay for itself at your size. The pilot settles that.

Backup

Why 56.8% is not accuracy

MEASURE	RESULT
Agreement with CFPB labels — gpt-4o-mini	56.8%
The same on gpt-4o, at 16.7× the cost	60.5%
Self-agreement, same input — gpt-4o	88.0%
Fabricated quotes caught	4 of 60

Consistent with itself, and the labels were picked by a complainant.

The anomaly detector's 100%

What it legitimately shows

Deterministic. All four patterns. 97.1% precision on 2,069 transactions.

What it does not show

Thresholds and data generator share an author — it measures self-agreement.

Same trap as the agreement figure. Only a pilot settles it.

AI Act classification

STEP

FINDING

1 • An AI system?

Yes, Art. 3(1)

2 • Prohibited under Art. 5?

No — one close call

4 • An Annex III use case?

No

6 • Art. 50 transparency

Drafts marked

Seven steps in full. Step 4 is the one that decides it.

What actually leaves the building

WHAT

WHERE IT GOES

The complaint narrative

To the model provider, in full

The quoted sentence

Stored in monitoring

The customer reference

Pseudonymised

Pseudonymisation protects the reference, not the content.

What the pilot has to prove

MEASURE

THRESHOLD

Team accuracy vs expert labels	$\geq 80\%$
Acceptance rate, per handler	75 – 97%
Ungrounded output reaching a handler	Zero
Correct abstention	$\geq 95\%$
Anomaly precision on disposition	$\geq 30\%$

Anomaly flagging

A detector decides what is unusual. The model only explains why.

☐	2026-06-20	2 d	threshold_structuring	150.7×	€4935.00	5	held	REJECT_VALUE_MISMATCH
☐	2026-06-20	2 d	amount_spike	54.8×	€995.00	1	proposed	OK_PROPOSED

Proposed

OK_PROPOSED

Amount spike

a single amount far above this account's own normal

Transactions	Total	Account median	Times normal
1	€990	€18	54.5×

Case note:

This candidate was raised due to a significant spike in the transaction amount, with a single transfer of EUR 990.0, which is far above the account's median transaction of EUR 18.15. The amount is approximately 54.5 times higher than what is typical for this account, indicating an unusual activity pattern.

Fastest check: Review the account's transaction history for any similar high-value transfers.

Nothing has happened yet. A person decides.

Escalated to an analyst

Dismissed — ordinary for this account

Needs more information

Checks

- ✓ input
- ✓ call
- ✓ parse
- ✓ scope
- ✓ confidence
- ✓ grounding

confidence 0.80 · 1377 ms · gpt-4o-mini · ref cx_65b1bc11dceaf3b

transfer · online_banking · IE