

AI CONSULTING CAPSTONE • ROUND 2

Assist

Three capabilities. One controlled decision process.

What works, what does not yet pay, and how to decide with evidence.

Round 1 set the direction. Round 2 built the evidence

KEPT

Mid-size EU financial services firm

Complaint triage remained the lead use case

The system proposes; a person decides

CHANGED

One implemented use case became three

A POC became a working MVP

Cost estimates became ROI, risk and compliance assessments

Decision now: a staged evidence programme — not deployment.

Three jobs share one accountability problem

Complaint triage

Routing queue · complaint evidence · handler decides

Reporting assistance

Report narrative · computed figures · named employee approves

Anomaly flagging

Flag explanation · transaction values · analyst decides

The task changes. The accountability problem does not.


One decision spine controls all three

Validate input → Model proposes → Guards verify → Human decides

In scope? Confident enough to propose? Every quote, figure or value present in the evidence?

The system never routes, files or blocks anything by itself.

The system proposes. A person decides.



One system, not three

the same checks and the same record serve all three — a fourth capability reuses them rather than starting again

Every step is kept

proposal, checks and human decision stored together — which is what makes agreement and override rates measurable at all


Minutes, not days

500 complaints classified in about a minute, eight at a time

The POC proves the mechanism – not accuracy

Proves: the workflow can propose, cite evidence and stop safely.

Does not prove: client performance or value.



The system works; the bespoke economics do not

36-MONTH ROI

Bespoke build • consultant oversight	-31.0%
Productised across five firms	+16.9%
Productised • oversight in-house	+96.6%

At 2,426 complaints a year, bespoke deployment breaks even in month 69.
Approximately 3,800 are needed.

€9,918 of €17,326 is cash — ombudsman fees and fraud losses, not saved time.

Three risks decide whether a pilot is defensible

Value assumptions fail on real data 4×5

Measure before committing to deployment. Fixed phase exits preserve the right to stop.

Handlers rubber-stamp proposals 4×4

Acceptance above 97% is a failure signal, measured per handler.

Oversight costs more than it saves 3×4

Measure effort during the shadow pilot from day one.

Not high-risk. Real-data use still has gates

AI Act: Triage and reporting are not high-risk.

Anomaly flagging requires counsel confirmation. AI literacy applies now.

GDPR: The full complaint narrative reaches the provider.

Close the DPA, transfer assessment and zero-retention controls before using real data.

The programme earns each next step

PHASE • WHAT IT BUYS

FEE • EXIT

0 • Ground truth, baselines, legal closure	€5,600 • go/no-go on pilot
1 • Configuration + 60-day shadow run	€12,600 • go/no-go on deployment
2 • Integration, training and handover	€14,000 • only if gates pass

€18,200 reaches the deployment decision. Nobody acts on a suggestion in shadow.

Hard gates: $\geq 80\%$ accuracy • 75–97% acceptance • zero ungrounded outputs.

The MVP makes every proposal accountable

One record connects the proposal, its evidence, every check and the human decision.

☐	2026-06-24	2 d	On track	Credit card	Disputes and fraud	Proposed	
--------------------------	------------	-----	----------	-------------	--------------------	----------	--

Proposed 15d left

Route to Payments and collections

Queue: Charged fees or interest you didn't expect

passed all checks

Because the customer wrote:

"I was charged a late fee while I was still trying to obtain the promised refund and dispute resolution."

> The full complaint

Checks

- ✓ input
- ✓ call
- ✓ parse
- ✓ scope
- ✓ confidence
- ✓ grounding

confidence 80% · 1759 ms · gpt-4o-mini · ref cx_896b75e4a1982a0c

Nothing has happened yet. A person decides.

If rerouting, the right team is

Account servicing

Note (optional)

why the proposal was wrong

Accepted — routed as proposed

Rerouted — send to a different team

Sent to a person for full review

Capstone Round 2 · Assist

10

Proceed to evidence — not deployment

Approve Phases 0 and 1 up to €18,200.

Close the legal and data-transfer gates before real-data processing.

Retain fixed phase exits and the stated pilot thresholds.

Do not fund the €14,000 deployment phase unless the evidence and commercial structure support it.

The POC has not yet proved client performance

MEASURE	RESULT
Complainant labels · gpt-4o-mini (deployed) / gpt-4o	56.8% / 60.5%
Self-agreement on the same input · gpt-4o	88.0%
Synthetic anomaly precision	97.9%
Accuracy / anomaly precision on client data	Not measured

Phase 0 creates expert labels; the shadow pilot tests real-world usefulness.

UC-1 and UC-3 are not high-risk; UC-2 needs counsel

STEP	FINDING
1 • AI system under Art. 3(1)	Yes
2 • Prohibited practice under Art. 5	None identified
4 • Annex III high-risk use	UC-1/3: no • UC-2: counsel
Articles 50 and 4	Marking + AI literacy

The intended use and human decision boundary drive the classification.

The full complaint narrative still leaves the bank

WHAT

WHERE IT GOES

Complaint identifier

No • replaced before the model call

Complaint narrative

Yes • sent to the provider in full

Evidence + pseudonymous reference

Stored in monitoring

Pseudonymisation protects the reference, not the content. Transfer controls are the gate.

Every pilot threshold must pass before deployment

MEASURE

THRESHOLD

Accuracy against expert labels	≥80%
Triage-time reduction	≥40%
Handler acceptance	75–97%
Oversight effort	≤3 days / quarter
Ungrounded outputs reaching a handler	Zero
Correct abstention	≥95%
Anomaly precision on disposition	≥30%