

AI CONSULTING CAPSTONE • ROUND 2

Assist

Three capabilities. One controlled decision process.

What works, what does not yet pay, and how to decide with evidence.

Round 1 set the direction. Round 2 built the evidence

KEPT

Mid-size EU financial services firm

Complaint triage remained the lead use case

The system proposes; a person decides

CHANGED

One implemented use case became three

A POC became a working MVP

Cost estimates became ROI, risk and compliance assessments

Decision now: a staged evidence programme — not deployment.

Three jobs share one accountability problem

Complaint triage

Routing queue · complaint evidence · handler decides

Reporting assistance

Report narrative · computed figures · named employee approves

Anomaly flagging

Flag explanation · transaction values · analyst decides

The task changes. The accountability problem does not.

One decision spine controls all three

Validate input → Model proposes → Guards verify → Human decides


In scope? Confident enough to propose? Every quote, figure or value present in the evidence?

The system never routes, files or blocks anything by itself.

The POC proves the mechanism — not accuracy

Proves: the workflow can propose, cite evidence and stop safely.

Does not prove: client performance or value.



The system works; the bespoke economics do not

36-MONTH ROI

Bespoke build • consultant oversight	-31.0%
Productised across five firms	+16.9%
Productised • oversight in-house	+96.6%

At 2,426 complaints a year, bespoke deployment breaks even in month 69.
Approximately 3,800 are needed.

Three risks decide whether a pilot is defensible

Value assumptions fail on real data 4×5

Measure before committing to deployment. Fixed phase exits preserve the right to stop.

Handlers rubber-stamp proposals 4×4

Acceptance above 97% is a failure signal, measured per handler.

Oversight costs more than it saves 3×4

Measure effort during the shadow pilot from day one.

Not high-risk. Real-data use still has gates

AI Act: Triage and reporting are not high-risk.

Anomaly flagging requires counsel confirmation. AI literacy applies now.

GDPR: The full complaint narrative reaches the provider.

Close the DPA, transfer assessment and zero-retention controls before using real data.

The programme earns each next step

PHASE • WHAT IT BUYS

FEE • EXIT

0 • Ground truth, baselines, legal closure	€5,600 • go/no-go on pilot
1 • Configuration + 60-day shadow run	€12,600 • go/no-go on deployment
2 • Integration, training and handover	€14,000 • only if gates pass

€18,200 reaches the deployment decision. Nobody acts on a suggestion in shadow.

Hard gates: $\geq 80\%$ accuracy • 75–97% acceptance • zero ungrounded outputs.

The MVP makes every proposal accountable

One record connects the proposal, its evidence, every check and the human decision.

☐	2026-06-22	3 d	12 d	on track	Credit card	Originations	0.80	proposed	OK_PROPOSED
☐	2026-06-17	8 d	7 d	on track	Checking or savings account	Account servicing	0.90	proposed	OK_PROPOSED
☐	2026-06-17	8 d	7 d	on track	Credit card	Credit reporting	0.80	proposed	OK_PROPOSED
☐	2026-06-16	9 d	6 d	on track	Checking or savings account	Disputes and fraud	0.80	proposed	OK_PROPOSED
☐	2026-06-14	11 d	4 d	due soon	Vehicle loan or lease	Asset recovery	0.90	proposed	OK_PROPOSED

Proposed 15d left OK_PROPOSED

Route to Disputes and fraud

Queue: Problem with a lender or other company charging your account

passed all checks

Because the customer wrote:

"Chase froze the funds, closed my account, and referred the matter to its fraud department."

> The full complaint

Nothing has happened yet. A person decides.

Accepted — routed as proposed

Rerouted — handler chose a different team

Sent to a person for full review

Checks

- ✓ input
- ✓ call
- ✓ parse
- ✓ scope
- ✓ confidence
- ✓ grounding

confidence 0.80 · 2769 ms · gpt-4o-mini · ref cx_db78f958aa4d3606

Proceed to evidence — not deployment

Approve Phases 0 and 1 up to €18,200.

Close the legal and data-transfer gates before real-data processing.

Retain fixed phase exits and the stated pilot thresholds.

Do not fund the €14,000 deployment phase unless the evidence and commercial structure support it.

The POC has not yet proved client performance

MEASURE	RESULT
Complainant labels · gpt-4o-mini (deployed) / gpt-4o	56.8% / 60.5%
Self-agreement on the same input · gpt-4o	88.0%
Synthetic anomaly precision	97.1%
Accuracy / anomaly precision on client data	Not measured

Phase 0 creates expert labels; the shadow pilot tests real-world usefulness.

UC-1 and UC-3 are not high-risk; UC-2 needs counsel

STEP	FINDING
1 • AI system under Art. 3(1)	Yes
2 • Prohibited practice under Art. 5	None identified
4 • Annex III high-risk use	UC-1/3: no • UC-2: counsel
Articles 50 and 4	Marking + AI literacy

The intended use and human decision boundary drive the classification.

The full complaint narrative still leaves the bank

WHAT

WHERE IT GOES

Complaint identifier

No • replaced before the model call

Complaint narrative

Yes • sent to the provider in full

Evidence + pseudonymous reference

Stored in monitoring

Pseudonymisation protects the reference, not the content. Transfer controls are the gate.

Every pilot threshold must pass before deployment

MEASURE	THRESHOLD
Accuracy against expert labels	≥80%
Triage-time reduction	≥40%
Handler acceptance	75–97%
Oversight effort	≤3 days / quarter
Ungrounded outputs reaching a handler	Zero
Correct abstention	≥95%
Anomaly precision on disposition	≥30%