

CAPSTONE ROUND 1 · PITCH TO CHLEO

Show her what the AI does

Complaint triage for a mid-size financial services firm

Ugo Ahukannah

Round 1 — research, dashboard, POC, monitoring, cost



A mid-size bank that does not trust AI

TAKEAWAY You never asked whether AI works. You asked what it is, and how you would see it.

THE FIRM

Financial services, mid-size

~500 staff, ~250,000 retail customers, EU. Current accounts, cards, personal and vehicle lending.

THE FEAR

“AI is not transparent”

Your words. You could not say what “the AI” would be, or how you would sign up.

THE PITCH

Three use cases

1. Complaint triage • 2. Anomaly flagging • 3. Reporting assistance. Ranked by how visible the reasoning is.

And two I refuse to build: a customer-facing chatbot, which fails in public in your customer's voice; and credit scoring, which is Annex III high-risk under the EU AI Act.




The dashboard: four numbers a CEO reads first




■ Why each one matters to you

- 16,839 complaints — the size of the job, not a sample
- 98.0% answered on time — the measure you already report to your regulator
- 12.7% cost money to resolve — the only figure here in your currency
- 46% sit in five of 64 categories — the opportunity, as a share not a promise
- Whose numbers? Not yours. Public regulator data standing in for your book until Phase 0. The shape transfers; the volume does not.

It runs, it is watched, and it stops itself



The workflow, running on one real complaint



The same decision, opened up in monitoring



60 decisions traced. ■ 52 proposed a team ■ 4 stopped, not confident enough ■ 4 stopped, quoted a sentence that was not there

Demo, not production. It picks the right team 60% of the time on public data — not good enough to route on its own. So it proposes, and a person decides. That is the honest configuration until a pilot on your data says otherwise.



€14,000 to a decision, not €24,500 to a deployment

Someone accountable for watching it — €5,600

Platform €1,800

€7,400 a year to run

▲ the AI itself — €0.29 a year

0

Discovery + labelling — €5,600

2 weeks. Two of your handlers label 300 complaints. Produces the ground truth you do not have.

1

Pilot, 60-day shadow run — €8,400

11 weeks. It proposes, nobody acts on it, we measure against your labels.

2

Deployment, only if it passes — €10,500

4 weeks. Assist-only. Autonomous routing is a separate decision.

Two assumptions: volume comes from the regulator's own complaint rate applied to a firm your size, so treat the euros as an order of magnitude. And the return rests on avoided ombudsman referrals — four a year covers the running cost. The pilot is what tests it.



What I most want your feedback on

I lead this pitch by admitting a weak number — 60% — instead of a confident one.

Does that earn a sceptical client's trust, or lose the room?



Backup — the three use cases, and the two refusals

TAKEAWAY Ranked by how visible the reasoning is, not by the size of the prize.

What I would build

- Complaint triage — reads a complaint, proposes a team, quotes its reason
- Anomaly flagging — every flag carries a reason, a human decides
- Reporting assistance — drafts what a person signs

What I refuse to build

- Customer-facing chatbot — fails in public, in your customer's voice
- Credit scoring — Annex III high-risk under the EU AI Act
- Predicting which complaints cost money — creates the wrong incentive



Backup — five categories are half of every complaint in the data

TAKEAWAY The model needs to be good at five things and honest about the rest.

Managing an account — 19.9%

Problem with a purchase on your statement — 11.2%

A lender or company charging your account — 5.6%

Fraud or scam — 4.8%

Closing an account — 4.6%

46.1%

of all complaints sit in just 5 of 64 categories.

Concentration like that is a property of complaint books, not of one firm.



Backup — volume and cost are not the same picture

TAKEAWAY Cards pay out 5.7× more often than vehicle loans. Routing has money attached.

Product	Complaints	Share	Closed with money paid
Checking or savings	5,835	34.7%	14.0%
Credit card	5,382	32.0%	17.2%
Money transfer	2,551	15.1%	10.6%
Vehicle loan	1,523	9.0%	3.0%

That ratio should hold for her, because it follows from how the products work — not from who the firm is.



Backup — every decision accounted for, including the refusals

TAKEAWAY 60 decisions traced. Nothing silently wrong.

WHAT HAPPENED	HOW OFTEN, IN 60 DECISIONS
Proposed a team, all checks passed	52
Quote was not actually in the complaint	4 — caught, sent to a human
Confidence below the 0.70 threshold	4 — caught, sent to a human
Silently wrong, with no record	0

A system that explains itself convincingly but falsely is more dangerous than one that says nothing.



Backup — why 60% is not the model's fault

TAKEAWAY The bottleneck is the labels, not the model. That is what Phase 0 buys.

What I measured

- 60.5% correct team (gpt-4o, 240 complaints)
- A 6× more expensive model buys 4 points
- So the bottleneck is not the model

Why the number is what it is

- Ask it the same complaint twice and it answers the same way ~9 times in 10
- So it is consistent — it just disagrees with the official label
- That label was picked by the customer filing the complaint, from a menu, untrained



Backup — three corrections made before any analysis

TAKEAWAY Each one would have produced a confident wrong answer.

WHAT THE DATA APPEARED TO SAY	WHAT WAS ACTUALLY HAPPENING
Complaints fell 73% in July	Publication lag — the regulator only publishes once a firm responds. Window cut at 2026-06-27
Handling time is 0 days	It measures the regulator's routing, not any firm's speed. Metric dropped
Timeliness 98.0% vs 0.6% untimely	Two different fields. Relabelled so a CEO is not shown a contradiction



Backup — the compliance position

TAKEAWAY A preliminary view with reasoning beats a confident label with none.

Question	Position for Round 1
EU AI Act tier	Limited risk — it classifies and routes, a human resolves. Not Annex III creditworthiness
What keeps it there	The human decision boundary. Autonomous routing is a new assessment, not an upgrade
GDPR	Complaint text is personal data. Legal basis, retention and a DPIA are Round 2
Open risk	Model hosting region — an EU customer's text must not leave the EEA by accident



Backup — every number, and where it came from

TAKEAWAY Nothing on these slides is an estimate dressed as a measurement.

Figure	Value	Source
Complaint corpus	16,839	CFPB public API, 2026-05-01 to 06-27
Complaint rate	3.7 per 1,000 accounts	FCA aggregate complaints 2025 H2
Ombudsman case fee	£650, 3 free	Financial Ombudsman Service 2025/26
Tokens per complaint	684 in / 49 out	Measured, n=12
Team accuracy	60.5%	gpt-4o, 240 complaints, seed 42