

CAPSTONE ROUND 1 · PITCH TO CHLEO

See exactly what the AI does

Complaint triage for a mid-size financial services firm

Ugo Ahukannah

Round 1 — research, dashboard, proof of concept, monitoring, cost



A mid-size bank needs AI it can inspect

TAKEAWAY You never asked whether AI works. You asked what it is, and how you would see it.

ASSUMED CLIENT PROFILE

Financial services, mid-size

~500 staff, ~250,000 retail customers, EU. Current accounts, cards, personal and vehicle lending.

THE FEAR

“AI is not transparent”

Your words. You could not say what “the AI” would be, or how you would sign up.

THE PITCH

3 use cases

1. Complaint triage • 2. Anomaly flagging • 3. Reporting assistance. Ranked by how visible the reasoning is.

And 2 I refuse to build: a customer-facing chatbot, which fails in public in your customer's voice; and creditworthiness scoring, an Annex III high-risk use under the EU AI Act.



4 numbers define the opportunity

Deploy ⋮

Complaint Operations

Mid-size financial services · 2026-05-01 to 2026-06-27 · 16,839 complaints across 567 firms

COMPLAINTS HANDLED

16,839

across 64 issue categories

ANSWERED ON TIME

98.0%

335 missed the deadline

COST MONEY TO RESOLVE

12.7%


19.7% needed some form of remedy

COULD BE AUTO-SORTED

46%


sit in just 5 of 64 categories

How many complaints arrive each week



Between 1,487 and 2,488 complaints a week, averaging 2,045, with no trend up or down across the period. Week to week it moves around — one quiet week sits a third below the busiest — but this is a continuous workload to be absorbed, not a spike to staff around.

Where complaints come from




Where they cost money



■ Why each one matters to you

- 16,839 public complaints — what I analysed, not your annual volume
- 98.0% timely responses — a CFPB field, not your performance
- 12.7% ended with monetary relief — not complaint-handling cost
- 46.1% fall in 5 of 64 categories — a hypothesis to test in your data
- These are CFPB figures across 567 firms. Phase 0 measures your own volume, mix, and outcomes.

Every decision is visible—and uncertain cases stop



The workflow, running on 1 public CFPB complaint

Inputs	Feedback	Input	Output
cx_70ce62144230423	auto_proposed: 0.00, evidence_verbatim: 0.00, guard: 0.00, team_correct: 1.00	customer_ref: cx_3a92b2805d197b6, narrative: I am filing a complaint against XXXX XXXX regarding improper account hand...	customer_ref: cx_3a92b2805d197b6, decision: HUMAN_REVIEW, environment: eval, evidence: I believe XXXX handled this situation improperly because: "I had already substant...
cx_6082fa3b198fbb1			model: gpt-4o-mini, product: Credit card, prompt_version: 0dede1f14c89
cx_31547828e0ec496			proposed_queue: Problem with a company's investigation into an existing problem
cx_ac67140a7eb558ff			proposed_team: HUMAN_REVIEW
cx_9bef635868448001			reason: evidence quote is not verbatim from the complaint
cx_3a92b2805d197b6			reason_code: REJECT_EVIDENCE_NOT_VERBATIM
cx_87e36abb0be22215			confidence: 0.8
cx_834ebc0aee3be2c0			evidence_is_verbatim: false
cx_27034e71ad78b82a			narrative_chars: 3124
cx_ef366880c3d33470			
cx_eefea43e9d32f2b0			
cx_9f6cc0a5a610bad6			
cx_b453eaab8baed4bc			
cx_8a23cbf4f63244b3			
cx_e2f22c7c58667edd			
cx_653b2733dc36c347			

The same demo case in LangSmith monitoring



60 evaluation cases traced. ■ 52 proposed a team ■ 4 stopped, not confident enough ■ 4 stopped, quoted a sentence that was not there

Demo, not production. Agreement with consumer-selected CFPB labels is 60.5% — not deployment accuracy. The demo stops for human review; Phase 1 runs in shadow, with no action on suggestions.



€14,000 to decide, not €24,500 to deploy

Consultant oversight: 8 days/year × €700 — €5,600

Platform: €150/mo × 12 — €1,800

Estimated annual operating cost: €7,400

▲ Estimated model calls — €0.29/year

0

Discovery + labelling — €5,600

2 weeks elapsed · 8 consulting days × €700. 2 of your handlers label 300 complaints; their time is excluded.

1

Pilot, 60-day shadow run — €8,400

11 weeks elapsed · 12 consulting days × €700. The system proposes; handlers do not act; results use the Phase 0 labels.

2

Deployment, only if it passes — €10,500

4 weeks elapsed · 15 consulting days × €700. Handlers may use proposals; autonomous routing needs separate approval.

2 assumptions: annual complaint volume uses a regulator benchmark for a firm your size; 4 avoided ombudsman referrals a year is the break-even hypothesis. Fixed phase fees are quoted; the pilot tests the benefit.



Supporting evidence

Appendix for questions and discussion



Why complaint triage comes first

TAKEAWAY Ranked by how visible the reasoning is, not by the size of the prize.

3 use cases assessed

- Complaint triage — proposed for Phase 0 and Phase 1
- Anomaly flagging — later option; your investigator reviews each flag
- Reporting assistance — later option; a named employee approves output

2 excluded uses + 1 guardrail

- Customer-facing chatbot — excluded; public errors use your firm's voice
- Creditworthiness scoring — excluded; Annex III high-risk
- Guardrail — predicted payout never determines routing priority



5 categories account for 46.1% of complaints

TAKEAWAY Public CFPB data suggests a narrow test; Phase 0 selects your categories.

Managing an account — 19.9%

Problem with a purchase on your statement — 11.2%

A lender or company charging your account — 5.6%

Fraud or scam — 4.8%

Closing an account — 4.6%

46.1%

of 16,839 public complaints sit in 5 of 64 categories.

Phase 0 tests whether your complaint book is similarly concentrated.



Complaint volume does not predict monetary relief

TAKEAWAY In public CFPB data, credit-card complaints ended with monetary relief 5.7× as often as vehicle-loan complaints.

Product	Complaints	Share	Closed with money paid
Checking or savings	5,835	34.7%	14.0%
Credit card	5,382	32.0%	17.2%
Money transfer	2,551	15.1%	10.6%
Vehicle loan	1,523	9.0%	3.0%

This is a public-data relationship; Phase 0 tests whether it appears in your complaint book.



All 60 evaluation decisions are accounted for

TAKEAWAY 60 evaluation cases; every outcome has a reason code.

WHAT HAPPENED	HOW OFTEN, IN 60 DECISIONS
Produced a proposal; validation checks passed	52
Quote was not actually in the complaint	4 — flagged for human review
Confidence below the 0.70 threshold	4 — flagged for human review
Cases without a recorded outcome	0

A system that explains itself convincingly but falsely is more dangerous than one that says nothing.



60.5% agreement makes expert labels essential

TAKEAWAY CFPB labels are not expert routing ground truth; Phase 0 creates it.

What I measured

- 60.5% agreement with CFPB-derived team labels (gpt-4o, n=240)
- A model costing 6× more improved agreement by 4 percentage points
- That comparison does not establish expert-routing accuracy

Why the number is what it is

- gpt-4o chose the same team on 89.0% of repeated runs
- It still agreed with CFPB-derived team labels only 60.5%
- Those labels came from consumer-selected issue menus, not expert routing



3 data corrections changed the story

TAKEAWAY Each one would have produced a confident wrong answer.

WHAT THE DATA APPEARED TO SAY	WHAT WAS ACTUALLY HAPPENING
Complaints fell 73% in July	CFPB publication lag — records appear after a company responds or enough time passes. Analysis ends 27 Jun 2026
Handling time is 0 days	It measures the regulator's routing, not any firm's speed. Metric dropped
Timeliness 98.0% vs 0.6% untimely	Two different fields. Relabelled so a CEO is not shown a contradiction



Human review supports a preliminary compliance view

TAKEAWAY This is a Round 1 working view, not a final legal classification.

Question	Round 1 working view
EU AI Act classification	Not identified as Annex III high-risk for this intended use; Round 2 legal validation is required
Decision boundary	The system proposes a route; a human confirms and resolves. Autonomous routing requires a new assessment
GDPR work	Round 2 assesses legal basis, retention, and whether a DPIA is required
Open risk	Hosting region and cross-border transfer controls are not yet verified



Every number is measured, sourced, or assumed

Measured, sourced, and assumed figures are labelled separately.

Figure	Value	Source
Complaint corpus	16,839	CFPB public API, 2026-05-01 to 06-27
Scenario complaint rate	3.7 per 1,000 accounts	FCA aggregate benchmark, 2025 H2
Ombudsman case fee	£650 after 3 free cases/year	Financial Ombudsman Service 2025/26
Tokens per complaint	684 input / 49 output tokens	Project measurement, n=12
Agreement with CFPB label	60.5%	gpt-4o, n=240, seed 42