CAPSTONE ROUND 1 · PITCH TO CHLEO

Show her what the AI does

Complaint triage for a mid-size financial services firm

Ugo Ahukannah
Round 1 — research, dashboard, POC, monitoring, cost

Public CFPB data · 16,839 complaints   All figures reproducible from the repo

Chleo's objection is trust, not capability

TAKEAWAY She never asked whether AI works. She asked what it is and how she would see it.

THE FIRM

Mid-size, EU

~500 staff, ~250,000 retail customers. Current accounts, cards, personal and vehicle lending.

THE FEAR

“AI is not transparent”

Her words at dinner. She could not say what “the AI” would be, or how she would sign up.

THE JOB

Make one decision visible

Not a platform. One place where she can watch the machine decide, and disagree with it.

Three use cases — and two I will not build

TAKEAWAY Ranked by how visible the reasoning is, not by the size of the prize.

What I would build

- Complaint triage — reads a complaint, proposes a team, quotes its reason
- Anomaly flagging — every flag carries a reason, a human decides
- Reporting assistance — drafts what a person signs

What I refuse to build

- Customer-facing chatbot — fails in public, in your customer's voice
- Credit scoring — Annex III high-risk under the EU AI Act
- Predicting which complaints cost money — creates the wrong incentive

Volume and cost are not the same picture

TAKEAWAY Where a complaint lands has money attached to it — cards pay out 5.7× more often than vehicle loans.

Product	Complaints	Share	Closed with money paid
Checking or savings	5,835	34.7%	14.0%
Credit card	5,382	32.0%	17.2%
Money transfer	2,551	15.1%	10.6%
Vehicle loan	1,523	9.0%	3.0%

Two thirds of volume sits in deposits and cards — and so does the money.

Five categories are half of everything that arrives

TAKEAWAY The model does not need to understand 64 things. It needs five, and it needs to say “I don't know” about the rest.

- Managing an account — 19.9%
- Problem with a purchase on your statement — 11.2%
- A lender or company charging your account — 5.6%
- Fraud or scam — 4.8%
- Closing an account — 4.6%

46.1%

of all complaints sit in just 5 of 64 categories.

That is what makes the scope small enough to be honest about.

You can watch it decide — and every branch ends at a person

TAKEAWAY Running now on the cohort n8n instance. Open any step and read the exact data going in and out.

Complaint arrives → Model proposes a team → Four guards check it → A person decides

The model's answer is a claim, not a result. It must return valid JSON, name a queue that exists for that product, clear 0.70 confidence, and quote the complaint verbatim. Any failure goes to a human, with the reason stated.

It caught itself fabricating a reason — four times

TAKEAWAY 60 decisions traced in LangSmith. Every one accounted for, including every refusal.

WHAT HAPPENED	HOW OFTEN, IN 60 DECISIONS
Proposed a team, all checks passed	52
Quote was not actually in the complaint	4 — caught, sent to a human
Confidence below the 0.70 threshold	4 — caught, sent to a human
Silently wrong, with no record	0

A system that explains itself convincingly but falsely is more dangerous than one that says nothing.

What it cannot do yet, and why that is not the model's fault

TAKEAWAY 60.5% team accuracy. Not deployable — and the reason changes what you should buy first.

What I measured

- 60.5% correct team (gpt-4o, 240 complaints)
- A 6× more expensive model buys 4 points
- So the bottleneck is not the model

Why the number is what it is

- The model agrees with itself 88–89% of the time — it is stable
- It disagrees with the public label systematically, not randomly
- That label was picked by the customer filing the complaint, from a menu, untrained

€ The model is not the cost. Oversight is.

TAKEAWAY 97% of the running cost is human review and platform. 0.004% is the model.

Running cost, per year	Amount	What it actually buys
LLM API	€0.29	684 input + 49 output tokens per complaint, measured
Platform and monitoring	€1,800	n8n hosting, LangSmith, logging
Quarterly accuracy review	€5,600	Someone accountable for watching it
Total	€7,400	Against a realistic labour saving of €4,578

Model choice is not a financial decision at your volume — it is an accuracy decision.

What I would do next, and what it costs

TAKEAWAY €14,000 commits you to a decision, not to a deployment.

0

1

2

Discovery + labelling — €5,600

2 weeks. Two of your handlers label 300 complaints. Produces the ground truth you do not have.

Pilot, 60-day shadow run — €8,400

11 weeks. It proposes, nobody acts on it, we measure against your labels.

Deployment, only if it passes — €10,500

4 weeks. Assist-only. Autonomous routing is a separate decision.

Break-even: the running cost is covered if this prevents four ombudsman referrals a year (€760 each). Whether it does is what the pilot is for — I am not claiming it does.

Stepping out of the pitch — what I want your feedback on

- Is “this is not deployable yet, and here is the number” the right thing to lead a client pitch with — or does it lose the room?
- Phase 0 sells a labelling exercise before any AI. Does that read as rigour, or as stalling?
- I refuse two use cases on stage. Does that build authority, or sound negative?

The one thing I most want challenged: *that an honest 60.5% is a stronger pitch than a confident 85%.*

Backup — three corrections made before any analysis

TAKEAWAY Each one would have produced a confident wrong answer.

WHAT THE DATA APPEARED TO SAY	WHAT WAS ACTUALLY HAPPENING
Complaints fell 73% in July	Publication lag — the CFPB only publishes once a firm responds. Window cut at 2026-06-27
Handling time is 0 days	It measures the regulator's routing, not any firm's speed. Metric dropped
Timeliness 98.0% vs 0.6% untimely	Two different fields. Relabelled so a CEO is not shown a contradiction

Backup — the compliance position

TAKEAWAY A preliminary view with reasoning beats a confident label with none.

Question	Position for Round 1
EU AI Act tier	Limited risk — it classifies and routes, a human resolves. Not Annex III creditworthiness
What keeps it there	The human decision boundary. Autonomous routing is a new assessment, not an upgrade
GDPR	Complaint text is personal data. Legal basis, retention and a DPIA are Round 2
Open risk	Model hosting region — an EU customer's text must not leave the EEA by accident

Backup — every number, and where it came from

TAKEAWAY Nothing on these slides is an estimate dressed as a measurement.

Figure	Value	Source
Complaint corpus	16,839	CFPB public API, 2026-05-01 to 06-27
Complaint rate	3.7 per 1,000 accounts	FCA aggregate complaints 2025 H2
Ombudsman case fee	£650, 3 free	Financial Ombudsman Service 2025/26
Tokens per complaint	684 in / 49 out	Measured, n=12
Team accuracy	60.5%	gpt-4o, 240 complaints, seed 42