

CAPSTONE ROUND 1 · PITCH TO CHLEO

Show her what the AI does

Complaint triage for a mid-size financial services firm

Ugo Ahukannah

Round 1 — research, dashboard, POC, monitoring, cost



Chleo's objection is trust, not capability

TAKEAWAY She never asked whether AI works. She asked what it is and how she would see it.

THE FIRM

Mid-size, EU

~500 staff, ~250,000 retail customers. Current accounts, cards, personal and vehicle lending.

THE FEAR

“AI is not transparent”

Her words at dinner. She could not say what “the AI” would be, or how she would sign up.

THE JOB

Make one decision visible

Not a platform. One place where she can watch the machine decide, and disagree with it.



Three use cases — and two I will not build

TAKEAWAY Ranked by how visible the reasoning is, not by the size of the prize.

What I would build

- Complaint triage — reads a complaint, proposes a team, quotes its reason
- Anomaly flagging — every flag carries a reason, a human decides
- Reporting assistance — drafts what a person signs

What I refuse to build

- Customer-facing chatbot — fails in public, in your customer's voice
- Credit scoring — Annex III high-risk under the EU AI Act
- Predicting which complaints cost money — creates the wrong incentive



This is the dashboard you would open



What you read first

- 16,839 complaints handled — the workload, not a sample
- 98.0% answered on time — the number you already report
- 12.7% cost money to resolve — the only figure in your currency
- 46% could be auto-sorted — the opportunity, stated as a share not a promise
- Every panel says what it means in plain English underneath. No statistics jargon anywhere on it.



Volume and cost are not the same picture

TAKEAWAY Where a complaint lands has money attached to it — cards pay out 5.7× more often than vehicle loans.

Product	Complaints	Share	Closed with money paid
Checking or savings	5,835	34.7%	14.0%
Credit card	5,382	32.0%	17.2%
Money transfer	2,551	15.1%	10.6%
Vehicle loan	1,523	9.0%	3.0%

Two thirds of volume sits in deposits and cards — and so does the money.



Five categories are half of everything that arrives

TAKEAWAY The model does not need to understand 64 things. It needs five, and it needs to say “I don't know” about the rest.

Managing an account — 19.9%

Problem with a purchase on your statement — 11.2%

A lender or company charging your account — 5.6%

Fraud or scam — 4.8%

Closing an account — 4.6%

46.1%

of all complaints sit in just 5 of 64 categories.

That is what makes the scope small enough to be honest about.



You can watch it decide — and every branch ends at a person


TAKEAWAY Running now on the cohort n8n instance. Open any step and read the exact data going in and out.



The model's answer is a claim, not a result. It must return valid JSON, name a queue that exists for that product, clear 0.70 confidence, and quote the complaint verbatim. Any failure goes to a human, with the reason stated.



Here it is, actually running



What you are looking at

- Green ticks: every node executed, on a real complaint from the corpus
- Bottom panel: the proposed team, the confidence, and the quote it relied on
- It routed to Disputes and fraud — which matched the complaint's actual label
- Nothing here is a mockup. You can click any node and read its real input and output.



It caught itself fabricating a reason — four times

TAKEAWAY 60 decisions traced in LangSmith. Every one accounted for, including every refusal.

WHAT HAPPENED	HOW OFTEN, IN 60 DECISIONS
Proposed a team, all checks passed	52
Quote was not actually in the complaint	4 — caught, sent to a human
Confidence below the 0.70 threshold	4 — caught, sent to a human
Silently wrong, with no record	0

A system that explains itself convincingly but falsely is more dangerous than one that says nothing.



One of those four, opened up

LangSmith

Application

- All applications
- Search
- Home
- Tracing
- Engine
- Monitoring
- Datasets & Experiments**
- Evaluators
- Annotation Queues
- Custom Apps
- Prompts
- Playground
- Studio
- Context Hub
- Deployments
- Sandboxes
- LLM Gateway

Personal / Datasets & Experiments

Legacy API usage detected

triage-round1-ebb7facb

Inputs

1	cx_70cef62f44230423
2	cx_6082fa3b198f1bb1
3	cx_31547828e0ecd496
4	cx_ac67140a7eb558ff
5	cx_9bef635868448001
6	cx_3a92b2805d197fb6
7	cx_87e36abb0be22215
8	cx_83debc0aee3be2c0
9	cx_27034e71ad78b82a
10	cx_ef366880c3d33470
11	cx_ee1ea43e9d32f2b0
12	cx_9fecc0a5a610bad6
13	cx_b453eaabd8aed4bc
14	cx_8a23cbf4f63244b3
15	cx_e2f22c7c58667edd
16	cx_653b2733dc36c347

Trace

Target

Feedback

Input

Output

Fields

customer_ref cx_3a92b2805d197fb6

narrative I am filing a complaint against XXXX XXXX XXXX regarding improper account hand...

product Credit card

Fields

customer_ref cx_3a92b2805d197fb6

decision HUMAN_REVIEW

environment eval

evidence I believe XXXX handled this situation improperly because : * I had already substant...

model gpt-4o-mini

product Credit card

prompt_version 0ded1e114c89

proposed_queue Problem with a company's investigation into an existing problem

proposed_team HUMAN_REVIEW

reason evidence quote is not verbatim from the complaint

reason_code REJECT_EVIDENCE_NOT_VERBATIM

confidence 0.8

evidence_is_verbatim false

narrative_chars 3124

Why this one was stopped

- The model was 0.8 confident and named a plausible queue
- But evidence_is_verbatim is false — its quote was not in the complaint
- reason_code: REJECT_EVIDENCE_NOT_VERBATIM. Routed to a person instead
- This is the record for every decision, including the ones where nothing happened.



What it cannot do yet, and why that is not the model's fault

TAKEAWAY 60.5% team accuracy. Not deployable — and the reason changes what you should buy first.

What I measured

- 60.5% correct team (gpt-4o, 240 complaints)
- A 6× more expensive model buys 4 points
- So the bottleneck is not the model

Why the number is what it is

- The model agrees with itself 88–89% of the time — it is stable
- It disagrees with the public label systematically, not randomly
- That label was picked by the customer filing the complaint, from a menu, untrained



The model is not the cost. Oversight is.

TAKEAWAY 97% of the running cost is human review and platform. 0.004% is the model.

Running cost, per year	Amount	What it actually buys
LLM API	€0.29	684 input + 49 output tokens per complaint, measured
Platform and monitoring	€1,800	n8n hosting, LangSmith, logging
Quarterly accuracy review	€5,600	Someone accountable for watching it
Total	€7,400	Against a realistic labour saving of €4,578

Model choice is not a financial decision at your volume — it is an accuracy decision.



What I would do next, and what it costs

TAKEAWAY €14,000 commits you to a decision, not to a deployment.

0

Discovery + labelling — €5,600

2 weeks. Two of your handlers label 300 complaints. Produces the ground truth you do not have.

1

Pilot, 60-day shadow run — €8,400

11 weeks. It proposes, nobody acts on it, we measure against your labels.

2

Deployment, only if it passes — €10,500

4 weeks. Assist-only. Autonomous routing is a separate decision.

Break-even: the running cost is covered if this prevents four ombudsman referrals a year (€760 each). Whether it does is what the pilot is for — I am not claiming it does.



Stepping out of the pitch — what I want your feedback on

- Is “this is not deployable yet, and here is the number” the right thing to lead a client pitch with — or does it lose the room?
- Phase 0 sells a labelling exercise before any AI. Does that read as rigour, or as stalling?
- I refuse two use cases on stage. Does that build authority, or sound negative?

The one thing I most want challenged: *that an honest 60.5% is a stronger pitch than a confident 85%.*



Backup — three corrections made before any analysis

TAKEAWAY Each one would have produced a confident wrong answer.

WHAT THE DATA APPEARED TO SAY	WHAT WAS ACTUALLY HAPPENING
Complaints fell 73% in July	Publication lag — the CFPB only publishes once a firm responds. Window cut at 2026-06-27
Handling time is 0 days	It measures the regulator's routing, not any firm's speed. Metric dropped
Timeliness 98.0% vs 0.6% untimely	Two different fields. Relabelled so a CEO is not shown a contradiction



Backup — the compliance position

TAKEAWAY A preliminary view with reasoning beats a confident label with none.

Question	Position for Round 1
EU AI Act tier	Limited risk — it classifies and routes, a human resolves. Not Annex III creditworthiness
What keeps it there	The human decision boundary. Autonomous routing is a new assessment, not an upgrade
GDPR	Complaint text is personal data. Legal basis, retention and a DPIA are Round 2
Open risk	Model hosting region — an EU customer's text must not leave the EEA by accident



Backup — every number, and where it came from

TAKEAWAY Nothing on these slides is an estimate dressed as a measurement.

Figure	Value	Source
Complaint corpus	16,839	CFPB public API, 2026-05-01 to 06-27
Complaint rate	3.7 per 1,000 accounts	FCA aggregate complaints 2025 H2
Ombudsman case fee	£650, 3 free	Financial Ombudsman Service 2025/26
Tokens per complaint	684 in / 49 out	Measured, n=12
Team accuracy	60.5%	gpt-4o, 240 complaints, seed 42